

Path-Parameter ILC with a Learned Correction Layer for High-Accuracy Industrial Robots

Pravin Oli

Erasmus Mundus Joint Master in Intelligent Field Robotics Systems (IFROS)

Universitat de Girona, Spain · Eötvös Loránd University, Hungary

pravin.oli.08@gmail.com

Abstract—The absolute accuracy of industrial robots is limited by unmodeled *drivetrain error*: joint elasticity, gear transmission error, friction, backlash and thermal drift. A recent path-parameter iterative learning controller (ILC) cancels this error but relies on an external laser tracker to learn and leaves several problems open. We re-implement that ILC faithfully on a simulated KUKA iiwa14 whose every axis is modeled as a two-mass flexible joint, and address three of those open problems: (i) the error is self-learned from an integrated *joint-side (output) encoder* rather than a laser tracker; (ii) a small learned model generalizes the correction to an *unseen* path with zero trials; and (iii) a temperature-aware layer compensates a non-repetitive thermal drift that a fixed ILC table cannot. The ILC reduces the Cartesian tool error by $\sim 32\times$ (to ≈ 1 mm), the learned layer recovers near-converged accuracy on a new path instantly, and the thermal layer matches an oracle re-learn. We are explicit about what the simulation does and does not establish.

1 Introduction

Industrial manipulators are highly *repeatable* but their *absolute* accuracy is limited by effects a rigid kinematic model ignores: joint elasticity, gear and cycloidal *transmission error*, Stribeck friction, backlash, and slow thermal drift [1, 3]. Iterative learning control (ILC) suits such *repeatable* errors because it improves a feedforward correction over repetitions of the same task [2].

Schwegel and Kugi [1] proposed a *path-parameter* ILC that indexes the learned correction by arc length $\lambda \in [0, 1]$ rather than time, so one table transfers across execution speeds; on a real 6-axis robot it reached $\approx 95\%$ error reduction in two trials. However, it uses an external *laser tracker* to learn, and its conclusion lists four open problems: speed-variation performance, contact tasks, a *model-based learning filter*, and *generalizing* the learned correction to different paths.

This paper re-implements that ILC and targets three open problems on a simulated, physically detailed KUKA iiwa14. Contributions: (1) the drivetrain error is self-learned from an *integrated joint-side encoder* instead of a laser tracker; (2) a small learned model predicts the correction for an *unseen* path with zero additional trials; and (3) a temperature-aware layer compensates a non-repetitive thermal drift that a frozen ILC table cannot.

2 Flexible-Joint Robot Model

We use the real iiwa14 (meshes, masses, inertias from the MuJoCo Menagerie [7]) in MuJoCo [6], and turn every hinge into a *two-mass flexible joint* [4] (Fig. 1): an actuated motor degree of freedom θ_m and a passive link degree of freedom θ_ℓ coupled by a torsional spring K and damper d . Per axis, K ranges 18000 $\rightarrow$ 3000 Nm/rad and d ranges 50 $\rightarrow$ 8 Nms/rad from base to wrist, giving structural natural frequencies $f_n \approx 10\text{--}195$ Hz. Each axis therefore obeys a coupled two-mass equation of motion,

$$J_m \ddot{\theta}_m = \tau_m - K\Delta - d\dot{\Delta}, \quad J_\ell \ddot{\theta}_\ell = K\Delta + d\dot{\Delta} - \tau_d, \quad (1)$$




Figure 1: Flexible-joint drivetrain model. Every axis of the real iiwa14 is a two-mass system: a motor inertia J_m and a link inertia J_ℓ joined by a torsional spring K and damper d . The controller plans on the rigid model ($\theta_\ell = \theta_m$); the link deflects, so $\Delta = \theta_m - \theta_\ell$ is a real transmission error. The link also receives a disturbance torque τ_d (transmission error, friction, backlash, thermal drift), unknown to the controller. The *joint-side / output encoder* reads the true link angle and feeds the ILC; the true Cartesian tool position is used only for evaluation.

where τ_m is the commanded motor torque, $\Delta = \theta_m - \theta_\ell$ is the elastic deflection, and τ_d is the link-side disturbance defined below. The free torsional mode of this pair sits at $f_n = \frac{1}{2\pi} \sqrt{K/J_\ell}$, which marks where structural *vibration*—rather than the repeatable, position-locked error the ILC targets—begins to dominate, and hence sets the frequency above which the learning filter must roll off.

The controller plans on the *rigid* model and never sees the springs, so the link lags the motor by an unknown deflection $\Delta = \theta_m - \theta_\ell$. On top of this, the link receives a disturbance torque

$$\tau_d = \underbrace{K \varepsilon(\theta)}_{\text{transmission}} + \tau_f(\dot{\theta}) + \tau_{bl}(\theta) + \underbrace{K k_T T(t)}_{\text{thermal}}, \quad (2)$$

with an angle-periodic gear error $\varepsilon(\theta) = \sum_k a_k \sin(n_k \theta + \phi_k)$ (0.3–0.8 mrad, $n_k = 13\text{--}37$ cyc/rev), Stribeck friction [5] $\tau_f = (F_c + (F_s - F_c)e^{-(\dot{\theta}/v_s)^2}) \text{sign}(\dot{\theta})$ ($F_s/F_c = 1.7$), a backlash dead-zone τ_{bl} , and a slow thermal offset whose temperature

$T(t)$ is heated by friction and is therefore *non-repetitive*. None of these is given to the controller; the ILC must discover the correction from the joint-side encoder alone.

3 Path-ILC and the Learned Correction Layer

Path-parameter ILC. Following [1], the correction is a table $u[\lambda]$ indexed by the path parameter. After trial $i-1$ the update at each index is the proportional-derivative law

$$\alpha_i[\lambda] = u_{i-1}[\lambda] + K_p e_q[\lambda] + K_d \dot{e}_q[\lambda], \quad (3)$$

where $e_q = q_r - \theta_\ell$ is the joint error from the output encoder. A Gaussian Q -filter smooths the update along the path,

$$u_i = Q * \alpha_i, \quad \sigma = \frac{\sqrt{\ln 2}}{2\pi f}, \quad (4)$$

so non-repetitive, high-frequency content (e.g. encoder noise) is not learned. We use the paper’s gains $K_p=1$, $K_d=0.01$ and cutoff $f=5$ Hz—an order of magnitude below the slowest structural mode of (1) ($f_n \approx 10$ Hz), so the table captures the repeatable, position-locked deflection while rejecting joint vibration and sensor noise. Indexing by λ rather than time decouples the correction from execution speed: the same table is replayed against arc length, so a slower or faster traversal reads out the identical feedforward. The motor command is $q_r + u[\lambda]$ (Fig. 2).

Learned correction layer. Classical ILC must re-run trials on *each* new path or operating point—impractical when the workpiece shifts slightly between cycles, or when the robot warms over a shift. We instead treat the converged correction as a function of *context* and learn it offline. A compact multilayer perceptron—two hidden layers of 256 units with ReLU activations (Fig. 3)—maps a low-dimensional descriptor of the desired path geometry to the full table $u[\lambda]$, and is trained by Adam on a mean-squared loss; for the thermal study a temperature-indexed linear model plays the same role. Crucially, every training target is an ILC table obtained the hard way—by actually running trials on the MuJoCo plant for that context—and the test context is held out, so the reported number is what the *predicted* feedforward achieves on the real plant for a path or thermal state the model never saw. The contribution is this architecture: a context-adaptive learned layer sitting on top of a classical, well-understood ILC, not a black box replacing it.

4 Experimental Results

All errors are Cartesian RMS at the tool tip, evaluated against the true forward-kinematic TCP, which the controller never uses for learning. Each path is sampled at $N=200$ arc-length points, every condition runs a handful of ILC trials, and all runs are seeded, so each figure is regenerated end to end by a single command. Table 1 summarizes the headline experiments (Fig. 4), under the realistic repeatable disturbances of (2).

Self-learning without a tracker. Using only the joint-side encoder, the ILC drives the tool error from 34 mm to ≈ 1.06 mm





Figure 2: Control structure. The path-ILC learns $u[\lambda]$ from the joint-side encoder error; the learned layer predicts $u[\lambda]$ for an unseen path or thermal state with zero trials. The true TCP is used only for evaluation.



Trained offline on ILC-converged tables collected for many path shapes; the test path is held out (zero trials at deployment). For thermal drift the same role is played by a temperature-indexed linear model.

Figure 3: Learned correction layer. A small multilayer perceptron maps the desired path geometry $\gamma(\lambda)$ to the entire feedforward table $u[\lambda]$ in one shot; it is trained offline on ILC-converged tables gathered for many path shapes, and the test path is held out so zero trials are run at deployment. The thermal variant replaces the network with a temperature-indexed linear model.

Table 1: Headline results (Cartesian RMS at the tool tip).

Experiment	Result
Convergence (output encoder only)	34156 $\rightarrow$ 1061 m ($\sim 32\times$)
Speed transfer (trained speed)	≈ 997 m (partial elsewhere)
Path transfer (no learned layer)	42016 $\rightarrow$ 14111 $\rightarrow$ 1663 m
NN, unseen path, 0 tr.	34594 $\rightarrow$ 5736 $\rightarrow$ 1461 m
Contact task (~ 108 N press)	18065 $\rightarrow$ 1055 m ($\sim 17\times$)
Learned, unseen thermal, 0 tr.	frozen 4508 $\rightarrow$ 988 $\approx$ oracle 1050 m

in 2–3 trials ($\sim 32\times$). A validation experiment (Fig. 6(b)) shows the encoder RMS and the true TCP RMS fall in lock-step (correlation 0.97), confirming the encoder is a faithful proxy for absolute accuracy.

Generalization to a new path. Reusing a table learned on path A barely helps on a different path B (42 $\rightarrow$ 14 mm); a full re-learn works but costs trials (1.66 mm). The learned layer predicts a correction for an unseen path and reaches 1.46 mm with *zero* trials.

Thermal drift. Thermal drift is non-repetitive, so a frozen table degrades as the robot warms (1.0 $\rightarrow$ 2.3 mm); an online ILC tracks it (Fig. 6(h)). A temperature-aware layer predicts the compensation at an unseen thermal state and reaches 0.99 mm, matching an oracle re-learn (1.05 mm).

Other KUKA experiments. Fig. 5 shows speed transfer




Figure 4: Headline results: (a) ILC convergence from the joint-side encoder only; (b) the neural-network layer on an unseen path, zero trials; (c) the temperature-aware learned layer at an unseen thermal state.

(best at the trained speed; partial elsewhere—the paper’s open speed-variation issue), path transfer without the learned layer, and a *contact* task in which the tool is dragged on a worktable under a ~ 108 N press: the same ILC reduces the in-plane error by $\sim 17\times$.

Detailed analysis. Fig. 6 collects the supporting figures: (a) the paper’s trial-wise plot and (b) the no-tracker validation; (c) the error along the path and (d) the spring–mass–damper view of the joints; and (e)–(h) an ablation isolating each realistic effect—the ILC stays ~ 1 mm under all of them, the Q-filter rejects encoder noise, and the injected transmission error and thermal drift behave as modeled.

5 Discussion: Honest Scope

This is a simulation, and two points deserve emphasis. First, the sensing assumption is *stronger* than the original hardware: our joint-side encoder reads the true link angle and thus *directly* observes the drivetrain error, whereas the laser tracker in [1] is needed precisely because motor-side encoders cannot. Real secondary encoders have noise and resolution limits, and a joint-to-task kinematic-calibration gap remains that a joint sensor cannot see. Second, speed transfer is only partial: a position-indexed table cannot cancel a velocity-dependent residual, an effect enlarged by our position controller relative to the paper’s computed-torque scheme. The learned layer is also an easy instance here (smooth low-dimensional paths); the contribution is the *architecture and protocol*—train on real ILC-solved tables, test on an unseen case—rather than the difficulty of this instance. Third, the gains are the paper’s defaults applied uniformly; a mild per-axis schedule (a higher K_p on the stiff wrist joints) would likely trim the residual further, so the absolute micrometre figures should be read as representative rather than optimized.

What would it take on hardware? The pieces map cleanly onto a real cell. The joint-side encoder is the secondary encoder that high-accuracy arms increasingly carry; the path-parameter table and Gaussian Q-filter are a few lines of feed-forward around the existing position loop; and the learned layer trains entirely offline, from tables the robot produces during its own ILC runs, so no extra instrumentation is needed at deployment. The risks are precisely the ones a simulation hides: finite encoder resolution and noise, the kinematic-calibration gap between joint space and task space that an output encoder cannot

observe, and load- or temperature-dependent dynamics richer than our two-mass model. That is why we frame the result as an *architecture* validated in simulation rather than a deployment claim—each open problem above is a concrete next experiment, not a hidden assumption.

Finally, a word on cost and division of labour. Nothing here is expensive at run time: the ILC update is one filtered PD pass per arc-length sample, and the learned layer is a single forward evaluation of a small network, both negligible beside the physics step. The design deliberately keeps the *repeatable*, structured part of the error with the classical ILC—where it is interpretable and provably convergent—and asks the network only to *interpolate between converged solutions* across contexts. That separation is why the learned layer can stay small and why a failure is diagnosable: if the prediction is poor on a new path, a few ordinary ILC trials recover it and simultaneously become fresh training data. The friction, transmission-error and thermal models we inject are standard ones [5, 3], so the difficulty the ILC faces is representative rather than adversarial.

6 Conclusion

We re-implemented a path-parameter ILC on a physically detailed flexible-joint KUKA iiwa14 [4] and extended it toward three of its stated open problems: tracker-free self-learning from a joint-side encoder, zero-trial generalization to new paths via a small neural network, and compensation of non-repetitive thermal drift with a temperature-indexed model. Across all of them the tool error stays near 1 mm, the limitations are stated plainly, and the methods are released as a single reproducible pipeline. The natural next step is hardware: a real secondary encoder, a kinematic-calibration stage to close the joint-to-task gap, and on-robot collection of the converged tables that train the learned layer. We expect the classical core to transfer directly and the learned layer to need richer context features (load, pose, temperature) than the smooth paths studied here demand.




Figure 5: Further KUKA experiments: (a) reusing the table at other speeds; (b) path A→B transfer without the learned layer; (c) ILC on a tool pressed on a worktable.




Figure 6: Analysis (a–d) and realistic-effect (e–h) figures. (a) trial-wise error / ILC inputs / path speed (trial 5 ILC-off, 6 on); (b) encoder RMS vs true TCP RMS in lock-step; (c) where the error lives along the path; (d) joint natural frequencies vs the Q-filter cutoff; (e) ILC stays ~1 mm under every effect; (f) Q-filter noise rejection; (g) the injected angle-periodic transmission error; (h) thermal drift defeats a frozen table while an online ILC tracks it.

References

- [1] M. Schwegel and A. Kugi, “A Simple Computationally Efficient Path ILC for Industrial Robotic Manipulators,” in *Proc. IEEE ICRA*, 2024, pp. 2133–2139. doi:10.1109/ICRA57147.2024.10610623
- [2] D. A. Bristow, M. Tharayil, and A. G. Alleyne, “A survey of iterative learning control,” *IEEE Control Syst. Mag.*, vol. 26, no. 3, pp. 96–114, 2006. doi:10.1109/MCS.2006.1636313
- [3] G. Litak and M. I. Friswell, “Vibration in gear systems,” *Chaos, Solitons & Fractals*, vol. 16, no. 5, pp. 795–800, 2003. doi:10.1016/S0960-0779(02)00452-6
- [4] M. W. Spong, “Modeling and control of elastic joint robots,” *ASME J. Dyn. Syst. Meas. Control*, vol. 109, no. 4, pp. 310–319, 1987. doi:10.1115/1.3143860
- [5] B. Armstrong-Hélouvry, P. Dupont, and C. Canudas de Wit, “A survey of models, analysis tools and compensation methods for the control of machines with friction,” *Automatica*, vol. 30, no. 7, pp. 1083–1138, 1994. doi:10.1016/0005-1098(94)90209-7
- [6] E. Todorov, T. Erez, and Y. Tassa, “MuJoCo: A physics engine for model-based control,” in *Proc. IEEE/RSJ IROS*, 2012, pp. 5026–5033. doi:10.1109/IROS.2012.6386109
- [7] Google DeepMind, “MuJoCo Menagerie,” 2022. [Online]. Available: https://github.com/google-deepmind/mujoco_menagerie